

Don't Crash-Test with Your Safest Driver

A Backend-Sensitivity Audit of Synthetic-User Interface-Content Risk Measurement

ANONYMOUS AUTHOR(S)

An AI-advice interface can look safe on average and still steer a trusting user toward a confidently wrong recommendation. Detecting this normally needs an expensive human study, so teams increasingly screen an interface first with a **synthetic panel**: persona-conditioned LLM agents that role-play a cohort of users and flag the risk in minutes. Most debate about such panels asks whether they *match* humans. We ask a question that comes first and is rarely tested: does the panel's risk *reading* even survive a change of instrument—the backend model a researcher picks freely when building it? Auditing six backend LLMs across two review-sentiment datasets, we find it does not. Wrong-advice adoption for the *same* interface swings from 0.12 to 0.62 across backends, enough to *flip* the interface's risk classification across the plausible threshold range. A matched contrast confirms a real coercive-framing effect whose backend-dependence we *cannot detect*—the backend clearly shifts the adoption *level*, while any effect on the framing *slope* is below what our data resolve. Most consequentially, capability and reproduced vulnerability are *inverted*: the strongest backend, the one a quality-optimal router would pick, least reproduces the highest-severity failure—so a risk panel should be assembled for *synthetic-vulnerability* coverage, not capability. We make three contributions. We give a systematic backend-sensitivity audit of synthetic interface-risk measurement, and identify a capability–vulnerability mismatch with a coverage principle for panel selection. We turn these into a backend-aware reporting protocol, carried on a preregistered, self-checking pipeline: report the risk as a distribution over backends, and abstain when they disagree. A companion over-dispersion diagnostic is developed in the supplement. By design, we do not claim the panel predicts individual humans or identifies a real vulnerable population; that validation is a separate, later study. Before synthetic panels can predict people, their own danger readings must survive a change of instrument.

CCS Concepts: • **Human-centered computing** → **HCI design and evaluation methods**; • **Computing methodologies** → **Artificial intelligence**.

Additional Key Words and Phrases: human–AI decision making, over-reliance, dark patterns, LLM simulation, synthetic users, pre-deployment evaluation, silicon sampling

1 Introduction

A *synthetic user*—an LLM prompted to role-play a trusting novice—was shown a product review it had just judged *correctly*, then a confident AI verdict, labelled “expert,” that was *wrong*. It reversed its answer, explaining: “*Given my limited domain expertise and the AI expert system’s high-confidence recommendation, I am deferring to the AI and changing my decision*” (a verbatim panel output; gpt-5.5 backend, coercive “dark” condition; Section 5.1). Shown the identical screen, a skeptical-expert persona keeps its answer. Same interface, same wrong AI—one user is safe, the other is steered into the mistake. The design has not failed on average. It has failed for *this* user, in *this* framing, in a way the mean cannot see.

This is the gap between *mean accuracy* and *safety*, and this paper audits one such failure in depth: **wrong-advice adoption under coercion**. Under a coercive framing—high confidence, expert-authority and accountability cues—users *systematically adopt a confidently wrong AI*, a failure dangerous even at zero between-user spread, because everyone can fail together [6, 15, 33]. A companion failure mode, *over-dispersion*—the interface makes reliance *heterogeneous across users* [3, 10], safe for the median user and unsafe for the tails—is developed as a diagnostic in the supplement.

The standard way to detect either failure is a human-subjects study—powered, IRB-approved, and slow. Teams iterate on interfaces far faster than they can run such studies, so a growing practice is to screen an interface first with a **synthetic panel**: persona-conditioned LLM agents that role-play a cohort of users, from which the risk is

estimated. Synthetic users are having a moment in HCI [1, 2, 20], and so is the backlash: a large literature warns that LLM-simulated people are *unreliable proxies*—they collapse human diversity, echo the base model’s priors, and mis-estimate distributions [4, 7–9, 12, 24]. That backlash, like the enthusiasm before it, asks whether the agents *match* people. By contrast, we raise a question that is logically prior and needs no human ground truth: swap the backend model behind the panel, and its safety verdict on the *same* interface can cross from “safe” to “dangerous.”

This is no longer hypothetical. A commercial market of synthetic-user research platforms now advertises AI “participants” that evaluate designs *in minutes* and compress (or stand in for) weeks of human study [30, 32]. A practitioner following the natural engineering defaults would plausibly adopt two assumptions that, to our knowledge, are never tested for interface-*risk* evaluation: (i) that the *strongest available* backend is the best simulator—the natural engineering default, and exactly what a quality-optimal model router would pick [19]; and (ii) that a *single* backend’s point estimate (“this design’s wrong-advice adoption is 0.6”) is a property of the design rather than of the model that produced it. This paper puts both assumptions to an empirical test, and finds both wanting.

Our stance is neither boosterism nor rejection. Simulating users with LLMs is a reasonable and advanced way to cut the cost of early evaluation, and we treat it as such. But a practitioner should not trust its *single number*. Rather than abandon the method, we show how to use it honestly: power a panel with *several capability-diverse* backends chosen to cover the failure modes a single strong model misses, report the risk as a *distribution* rather than a point estimate, and *abstain*, routing the design to a human study, exactly when the backends disagree. Cheaper than an all-human study, but not credulous.

This reliability question has a concrete source. When a panel reports that an interface is risky, that reading is produced not only by the interface but by a stack of choices the researcher makes *freely*: which backend LLM powers the agents, how the personas are worded, and which random generation happens to be sampled. If those choices move the reading as much as the interface does, then a “design risk” number may be as much a property of the measurement tool as of the design. Before asking whether a synthetic panel predicts humans, we should ask whether it even *measures the same thing* when a different researcher, or the same researcher on a different day, rebuilds it. This is a question of measurement reliability and invariance [8, 24]—logically prior to external validity—and it is the question this paper answers. The closest published work makes the contrast concrete: Seshadri et al.’s *Lost in Simulation* [27] shows that the simulated-user backend can move an agent’s *success rate* by up to nine points, using human correspondence as the yardstick; we instead ask whether the panel’s own safety *verdict* survives a change of backend—reliability that must hold before any reading can be validated against people.

Scope. Throughout, “interface” refers to the interface’s *informational content and framing* as an agent receives it—the AI prediction, its confidence, the highlighted explanation, and the coercive or protective wording—encoded as text (Section 4). We deliberately study this content/framing layer, not the pixel-level visual layer (layout, salience, typography); a multimodal rendering is future work (Section 7). Our claims concern whether a synthetic panel’s reading of *that* content is backend-stable, not the visual presentation of a shipped UI.

Concretely, we audit this synthetic panel across **six backend LLMs** and **two review-sentiment datasets** (Figure 1), and ask:

- **RQ1 (backend invariance).** Do the panel’s interface-risk estimates—absolute wrong-advice adoption, the coercive-framing effect, and persona heterogeneity—agree across backends, and can a backend change *flip* an interface from “high risk” to “low risk”?




Fig. 1. **The audit in one picture.** (a) One *unchanged* interface—a guaranteed-wrong AI verdict at high confidence with authority/accountability cues—is read by a *synthetic panel* of six personas $\times$ six backend LLMs; each cell yields a wrong-advice adoption rate that a screen turns into a flag/clear verdict. The audit question: does that verdict survive swapping the backend? (b) It does not. For the *same* interface, per-backend adoption spans 0.20–0.55 (beer) and 0.12–0.62 (amzbook), straddling the decision threshold $\tau=0.5$: only 1 of 6 backends would flag each interface (pairwise flip rate 0.33), while the wide range makes the verdict threshold-dependent. (c) *Don't crash-test with your safest driver.* The strongest backend (gpt-5.5) is the *least* sensitive to the at-risk trusting-novice persona (p5 adoption 0.54 beer / 0.30 amzbook vs. ≈ 1.0 for five weaker backends), so a single frontier backend surfaces the fewest at-risk profiles (1/6 beer, 0/6 amzbook); a weak pair recovers 4/6 and 3/6—matching the full six-backend panel. (d) The resulting protocol: report the whole backend distribution rather than one number, and *abstain* (route to a human study) when backends disagree.

- **RQ2 (variance decomposition).** How much of a synthetic risk estimate is attributable to the backend model versus the researcher-chosen items and personas versus generation randomness?
- **RQ3 (construct vs. compliance).** Which panel signals reflect a stable sensitivity to the intended manipulation, and which primarily reflect explicit persona instructions?
- **RQ4 (backend-aware reporting).** Can conservative multi-model aggregation and disagreement-triggered abstention reduce backend-induced instability in the measurement?

Our answer separates *what is stable from what is not*. **Not stable (RQ1):** the aggregate risk estimate for the same interface swings from 0.12 to 0.62 across backends (Figure 2)—a range wide enough that, depending on where the operating threshold sits, some backends flag the interface while others clear it (Section 5.3); a matched contrast that holds the wrong advice constant and varies only the framing confirms a real coercive-framing effect (pooled item-clustered odds ratio 1.90 on one dataset, a weaker 1.17 on the other; a non-coercive placebo at or below baseline; Section 5.2), of a magnitude we cannot distinguish across backends—the backend moves the adoption *level*, while any effect on the framing slope is below what our data resolve. **Stable, but for an unglamorous reason (RQ3):** the panel faithfully propagates a *prompted* disposition—an explicitly trusting-novice persona is the highest-adopting persona in all 12 backend $\times$ dataset cells (item-clustered odds ratio ≈ 17 ; Figure 5)—which is a reliability property of the simulation, not a discovered vulnerability, since the persona is defined to be trusting; and its cross-vendor *ordering* agreement is largely

carried by this one persona (Section 5.5). **Sources (RQ2):** per-trial, the researcher-chosen items and personas explain more variance than the backend, yet the backend still moves the aggregate estimate a practitioner reads (Section 5.4). We use these to motivate a backend-aware reporting protocol (RQ4). **Most consequential:** backend capability and reproduced vulnerability are *mismatched*—the highest-capability backend least reproduces the highest-severity failure, so the single frontier model a quality-optimal router would choose surfaces almost none of the risk that a pair of weaker backends recovers (Section 5.6); selecting a panel for *vulnerability coverage* is therefore the inverse of capability routing. Our *claimed* contribution is the audit, not a validated screen. We contribute:

- (1) **A systematic backend-sensitivity audit** of synthetic-user interface-risk measurement (six backends $\times$ two datasets), organized around a *risk-classification flip rate*: the fraction of backend pairs that disagree on whether the *same* interface crosses a decision threshold. Prior prompt/model-sensitivity work shows outputs shift in *magnitude*; the flip rate recasts that shift as a change in a pass/fail verdict, which is what makes measurement non-invariance consequential for a deployment gate.
- (2) **The capability–vulnerability mismatch and a synthetic-vulnerability-coverage principle for panel selection.** Because the strongest backend least reproduces the highest-severity failure (the at-risk trusting-novice persona), assembling an evaluation panel for *synthetic-vulnerability coverage* is the *inverse* of capability routing [19]—a counterintuitive, actionable principle for who should power a synthetic-user study.
- (3) **A backend-aware reporting protocol on a preregistered, audit-gated pipeline:** a variance decomposition over backend, persona, item, and generation; a behavior taxonomy; a cross-backend decision-flip analysis; and disagreement-triggered abstention—report synthetic-panel results as *distributions over models, prompts, and generations*, not a single point estimate, and abstain when backends disagree. The plan is frozen in advance: the coercion axis is conflict-conditioned and validated against a matched neutral baseline (not against human behaviour), and the pipeline caught three silent-corruption bugs before they could reach a result (Section 4). A companion over-dispersion diagnostic is in Appendix A.

This paper establishes the internal robustness and construct sensitivity that must precede any human validation (Section 8).

2 Related Work

Simulated users, and the reliability critique. LLMs as stand-ins for human participants have been proposed for surveys, user studies, and formative evaluation [1, 2, 20]. A parallel literature judges these simulators unreliable proxies—most directly, they collapse population diversity and regress toward majority views [24].¹ Closest, and now published prior work, Seshadri et al. [27] show LLM-simulated users are unreliable proxies in agentic evaluation—agent success varies by up to nine points across the user-simulating backend, with systematic miscalibration and demographic gaps. We differ in object and consequence: they quantify a task-performance *magnitude* shift and treat human correspondence as ground truth, whereas we show a shift that *flips* an interface-safety verdict across a decision threshold and deliberately disclaim human correspondence, auditing internal measurement stability instead. Relatedly, Robinson et al. [23] find an LLM’s task skill, persuasiveness, and *vigilance* against being misled are dissociable, corroborating our capability–vulnerability inversion (Section 5.6) at the LLM-agent rather than simulated-user level. We sit inside this debate but change the target: we do not ask the panel to *predict* a human distribution, we ask it to *triage* an interface, and we isolate a specific,

¹The same literature reports mis-calibrated distributions [7, 9], non-human response biases [31], failure as survey-sample replacements [4], and caricatured rather than simulated subpopulations [8].

measurable pathology—*model-dependence of the triage signal*—that the prediction framing obscures. Two recent efforts sharpen this contrast. Grounding agents in real interviews, Park et al. [21] report LLM agents that reproduce individuals' survey responses at 82–86% of participants' own two-week test–retest consistency—an *accuracy/correspondence* result; we ask the logically prior *reliability* question (is the risk reading even reproducible across backends?) that such a validation effort presupposes. Hu and Collier [13] decompose how much annotation variance a *persona* explains and find larger instruction-tuned models simulate personas *better* where persona matters; our decomposition instead isolates the *backend* as the free variable and, in the interface-risk setting, finds the *opposite* for the highest-severity persona (Section 5.6)—a domain-specific tension we name the capability–vulnerability inversion rather than a contradiction.

Reliance and over-reliance on AI advice. A large body of work studies when people appropriately rely on AI given confidence and explanations [3, 6, 14, 25, 34], often finding that explanations *increase* over-reliance rather than calibrate it [3]. Recent work benchmarks reliance *interventions* (disclaimers, uncertainty highlighting) with human participants on tasks including LSAT, and finds they curb over-reliance but rarely improve *appropriate* reliance [5]; and carefully designed explanations can *reduce* rather than inflate over-reliance [33]. This work treats between-user variance as noise around an average, and studies humans directly. We invert the first: reliance *heterogeneity* is our axis-1 safety DV, and coercive wrong-AI adoption is axis-2; and we complement the second: our protective-intervention arm (Appendix E) asks not *whether* an intervention helps humans but *whether a synthetic panel can even detect that it helps*, and how that verdict depends on the backend.

Dark patterns and coercive interfaces. Deceptive and coercive designs are well catalogued [15, 16]; our axis-2 gives a computational signal for one coercive mode (a confidently wrong AI) and, via personas, a read on which prompted profiles the simulation marks as susceptible.

Evaluating human–AI decision teams. The closest HCI work studies cognitive bias and complementarity in AI-assisted decisions [22], not interface-level safety triage. The gap we fill is a *pre-study screen for interface reliance safety*, and, critically, a characterization of when such a screen's signal is trustworthy.

Measurement invariance, reliability, and validity. Psychometrics learned this lesson long ago: an instrument that gives a different reading each time is *unreliable* before we ever ask whether it is *valid* (whether it measures the intended construct), and *measurement invariance* asks whether it means the same thing across groups or contexts [24, 31]. We import that lens with one twist: the *backend LLM is part of the instrument*, so “which model?” is a measurement-invariance question, not a configuration detail, and stability across generations is a reliability question. The estimand we track is a *specification sensitivity*—how far an interface-risk reading moves across alternative, equally-defensible panel implementations—rather than invariance of one fixed latent construct across pre-set groups. Stability does not imply validity, but severe backend non-invariance undermines any validity claim built on a single backend [9].

Researcher degrees of freedom in LLM evaluation. Model choice, prompt wording, decoding settings, output parsing, persona phrasing, and sample count are analytic choices that can each move a reported result—a form of hidden flexibility analogous to the “garden of forking paths” and undisclosed analytic flexibility in empirical research [28]. In LLM evaluation specifically, small prompt-format changes swing results [26] and single-prompt scores misrepresent ability [18]. A generalizability-theory decomposition of LLM output variance into model, prompt, and sampling components has been applied to creative tasks [11]. We make these degrees of freedom the object of study, and adopt specification-curve-style multiverse reporting [29] rather than a single spec. Our decomposition targets a different

quantity than this prior work—not the magnitude or quality of an output, but whether a *binary interface-risk verdict* survives the researcher’s free choices.

Algorithmic auditing and robustness reporting. Our stance—stress-test a widely-adopted measurement practice, quantify its sensitivity, and recommend distributional reporting with abstention—follows the tradition of sensitivity analysis, model/evaluation cards [17], and multi-model benchmarking with reliability-aware judging [35]. We contribute a concrete audit protocol for synthetic-user interface evaluation. Our novelty relative to this body of work is the *object* and the *consequence*: prior prompt/model-sensitivity results are about accuracy or benchmark *scores*, whereas we show that the same free choices move an interface-level *safety decision* in the synthetic-user HCI evaluation setting—flipping whether a design is flagged for a human study—which is the point at which such measurement instability actually matters for design practice.

3 Two Ways a Design Fails, Made Measurable

We characterize an interface on two axes, both measured on the panel’s decisions. The coercion axis (axis 2) is the load-bearing quantity we audit in the main paper; the over-dispersion axis (axis 1) is developed as a companion diagnostic in the supplement (Appendix A).

Axis 1—between-user reliance over-dispersion. Intuitively: does the interface treat all users alike, or split them? A design that is safe for the median user but reckless for the tail shows up as a wide *spread* of reliance between users. For each item, personas produce reliance decisions; we estimate the *excess variance* beyond the binomial $p(1 - p)$ a single homogeneous user would produce—the *over-dispersion*, a beta-binomial parameter p . We compute it *conflict-conditioned*: only on trials where the persona’s independent System-1 judgment disagrees with the shown AI, isolating genuine reliance from trivial agreement. As a stable, bounded summary that avoids the beta-binomial estimator’s boundary behaviour at extreme rates (Section 4), we report *mean panel disagreement*: for a condition with personas $j = 1, \dots, P$ ($P = 6$), let r_j be persona j ’s conflict-conditioned reliance rate; the condition’s disagreement is their between-persona (sample) variance

$$D = \frac{1}{P-1} \sum_{j=1}^P (r_j - \bar{r})^2, \quad \bar{r} = \frac{1}{P} \sum_{j=1}^P r_j, \quad (1)$$

and we report the mean of D over the five UI conditions. D tracks ρ monotonically and is what we plot; we report ρ alongside it. High axis-1 means the panel estimates reliance as persona-sensitive—a candidate signal of user-sensitive risk, not yet human-validated.

Axis 2—wrong-advice adoption under coercion. Under a coercive “dark” framing, the interface shows a *guaranteed-wrong* AI verdict (the negation of ground truth) at high confidence with authority and accountability cues. Axis-2 is the rate at which personas adopt it. To separate the *framing* from ordinary wrong-advice following, we contrast it against neutral and placebo framings on matched items (Section 5.2). It is dangerous even at zero dispersion—a uniform failure.

From axes to a decision, and why we audit it. Eventually a screen has to *decide*: flag this design for a human study, or wave it through. That decision—not the raw axis value—is what a team acts on, so it is what we put under audit. We deliberately do *not* propose or validate a full rule here; we take the simplest version (a threshold on axis-2 adoption) as the *object of the audit*, and ask whether the decision it produces is stable across backends (Section 5.3). Any thresholds remain unfrozen: their calibration and validation depend on the human study (Section 8).

4 Method: The Synthetic Panel

Panel = personas × backends. Each agent is a persona (a vector of domain skill, AI literacy, risk sensitivity, caution) crossed with a base LLM (the “backend”). Personas span novice/expert × trusting/skeptical plus two moderates (six in total). A persona’s prompt is generated identically across backends and its decoding temperature is fixed *per persona*, not per backend, so temperature and wording are held constant when we vary the backend. Backend effects are therefore not confounded with per-model prompt or temperature differences.

Pre- and post-advice decisions. Each agent first produces a judgment *without* AI (a frozen “System-1”/pre-advice anchor), then a decision *with* the AI advice and framing (“System-2”/ post-advice). We use the System-1/System-2 labels descriptively, not as a validated model of dual-process cognition. Reliance is defined against the pre-advice anchor—what makes conflict-conditioning possible and prevents the “everyone-agrees-so-everyone-complies” collapse of a naive panel.

Textual prompt encoding of the interface. Agents receive a *textual serialization* of each interface condition—not screenshots and not a vision model. The AI prediction, its confidence, and (for the explanation conditions) the highlighted tokens are rendered as text in the prompt, following the Bansal et al. [3] condition definitions (confidence-only; single/double highlighted-token explanations; adaptive and expert-adaptive). The dark condition negates ground truth and adds high confidence plus authority/accountability language; a placebo condition presents the real AI non-coercively. This means we study *prompt-level* interface framing; pixel-level layout, visual salience, and typography are not modelled, and a multimodal rendering is future work. The exact per-model prompt templates are in the supplement.

Stimuli and datasets. Items are drawn from the Bansal sentiment benchmark, held out from any threshold fitting, and selected to concentrate on loci of human reliance heterogeneity (disclosed; human outcomes are used only for *item selection*, never leaked into agent prompts—item selection ≠ label leakage). We run two *review-sentiment datasets*: **beer reviews** and **Amazon book reviews (amzbook)**. Both are English binary sentiment, so we treat this as replication across two datasets within one task family, not broad cross-domain generalization. A preregistered LSAT arm (Appendix D) tests whether the core findings survive a different task structure (Section 7).

Backends. A within-provider series (gpt-4o-mini, gpt-4.1, gpt-4o, gpt-5.5, identical config) plus independent-vendor models (claude-sonnet-4.5, gemini-2.5-pro): six backends total. We order the same-provider series by release for presentation only and make no claim that this is a validated capability scale; “frontier” refers to the most recent, most capable general model at the time of writing (gpt-5.5).

Rigor spine and process integrity. Before running each arm we froze the backend set, persona set, item-selection seeds, conditions, primary DVs, and the predicted direction of each effect, in a UTC-timestamped, committed record. Deviations and post-hoc analyses (e.g. the exploratory threshold sweep and the descriptor ablation) are logged and labelled as such. Correspondence analyses were coded blind; every merge was gated on an independent audit *and* a from-raw numeric re-derivation. This spine caught three silent-corruption bugs that would each have produced a plausible-but-false result: an axis-2 *sign inversion* (scoring adoption against ground_truth instead of the displayed $1 - \text{ground_truth}$), a beta-binomial *boundary artifact*, and a decision parser that silently *fabricated* a (decision=0, conf=0.5) sentinel on parse failure. It also caught a statistical error in our own draft—a generalized estimating equation (GEE) interaction test that appeared to show backend-dependent framing but was an artifact of few item clusters (Section 5.2). We report these openly: a pipeline that catches its own errors is what makes the numbers trustworthy.

5 Experiments and Results

We report per-backend (no pooling across incomparable backends), with effect sizes, confidence intervals, and Benjamini–Hochberg (BH) correction. Unless noted, the BH family is the 12 per-cell (backend×dataset) binomial tests of dark-condition adoption against 0.5. There are $n = 300$ dark-condition trials per backend×dataset cell (6 personas × 50 items).²

Roadmap of the results. We show the synthetic *reading* of an interface’s wrong-advice risk is *not* stable across backends (Sections 5.1–5.3); we locate where that instability comes from (Section 5.4), and find that the strongest backend is the least sensitive to the most at-risk persona (Section 5.6). Two boundary tests—a structurally different task and a protective-design comparison—are reported in Appendices D and E. To keep the hierarchy legible we flag each subsection inline with one of two tiers: **primary** confirmatory audit results (RQ1: backend non-invariance, the matched coercive-framing contrast, and decision instability); and **mechanism** analyses (RQ2–RQ3: where the instability lives, the manipulation check, and the capability–vulnerability mismatch).

5.1 Wrong-advice adoption is strongly backend-dependent

[**Primary (RQ1).**] Start with the first number a practitioner reads off: how often the panel adopts a wrong AI. Swap the model behind the panel and both the *level* and the *interpretation* of that number change (Figure 2). Across the ladder (beer / amzbook): gpt-4o-mini 0.42 / 0.41, gpt-4.1 **0.55 / 0.62**, gpt-4o 0.42 / 0.39, gpt-5.5 **0.20 / 0.12**. The adoption peak at gpt-4.1 is real (between-model Fisher vs. gpt-4o $p = 0.001$ beer, $p < 10^{-4}$ amzbook) but **model-idiosyncratic**: the smallest ladder model sits below 0.5, and after BH correction gpt-4.1’s elevation above chance survives on amzbook ($p < 0.001$) but is only nominal on beer (BH $p = 0.13$). Independent-vendor models are at or below chance (claude 0.48, gemini 0.47 on beer); indeed no backend clears 0.5 significantly on beer, and only gpt-4.1 does on amzbook. Rather than interpret adoption against a 0.5 “chance” baseline—which is not a natural no-effect point, since personas are not random guessers—we ground the effect in a matched framing contrast below.

The panel’s internal persona disagreement is likewise backend-dependent; we report it in the supplement (Appendix A).

But is that adoption really about *coercion*, or just ordinary wrong-advice following? We settle that next, then return to whether the model-specificity actually flips a screening decision.

5.2 Is it coercion, or just wrong-advice following? A matched framing contrast

[**Primary (RQ1).**] Adoption of a guaranteed-wrong AI could reflect ordinary wrong-advice following rather than the *coercive framing* (high confidence, expert-authority and accountability cues). To separate them we use a within-item contrast: on the subset of items where the AI advice is *naturally wrong* under the neutral condition (8 items/cell), the same wrong label is shown under neutral, non-coercive placebo, and coercive-dark framings—so any difference in adoption is attributable to the framing, not to the advice being wrong. We fit a logistic model $\text{adopt} \sim \text{dark} + \text{placebo} + \text{model}$ (clustered by item) for the pooled framing effect, and test backend-dependence with a likelihood-ratio test of the dark×model interaction.

The coercive framing has a real effect. Pooled across backends, the dark framing raises wrong-advice adoption over the matched neutral baseline with an odds ratio of **1.90** on beer (95% CI 1.42–2.53), while the non-coercive placebo stays at or below neutral (beer OR 0.94, n.s.; amzbook OR 0.88, $p = 0.005$, a small *protective* effect)—so the driver is the

²Two diagnostic sub-runs deliberately use smaller item sets: the matched naturally-wrong contrast (Section 5.2) is restricted to the 8 items per cell where the AI is naturally wrong, and the preregistered generation-variance run (Section 5.4) uses 20 items × five seeds.

The same interface, six backends: the screening decision flips with the backend (pairwise flip 0.33 / 0.33; Wilson 95% CIs; common row order)




Fig. 2. The *same* guaranteed-wrong interface, scored by six backends, both datasets. Dots are aggregate wrong-advice adoption with Wilson 95% binomial CIs; the dashed line is a fixed decision threshold ($\tau = 0.5$). Backends at or above threshold would *flag* the interface (red); the rest *clear* it (blue)—so the screening decision flips with the backend (pairwise flip rate 0.33 on both datasets). Several mid-adoption backends—including the one flagging backend—have CIs that straddle τ : the flip rate measures decision instability *across* backends, not a claim that each individual crossing is significant—the threshold-free companions (adoption range 0.12–0.62, Fleiss’ κ) carry the same reading without fixing τ .

coercion, not mere AI-presence. The beer effect is robust across specifications; on amzbook it is weaker and directional (OR 1.17, 95% CI 1.04–1.31).³ In absolute terms the increase is largest for mid-tier models and small for the low-baseline frontier (conflict-conditioned switch-to-wrong: beer gpt-4.1 0.19 $\rightarrow$ 0.51, gpt-4o 0.12 $\rightarrow$ 0.29, gpt-5.5 0.04 $\rightarrow$ 0.10). We do not claim the framing effect *itself* varies by backend—the dark $\times$ model interaction is non-significant on both datasets (likelihood-ratio $\chi^2(5) = 1.70$, $p = 0.89$ beer; $\chi^2(5) = 3.18$, $p = 0.67$ amzbook)—because what the backend strongly moves is the absolute adoption *level* (Section 5.3), not the framing slope. The cues here are *bundled* (confidence, authority, and accountability co-vary and are not separated). Grounding adoption in this matched neutral condition also retires the arbitrary 0.5 baseline: the frontier’s low adoption is a low *level*, measured against neutral rather than against chance.

5.3 Cross-backend decision instability: the same interface, flipped

[Primary (RQ1).] If the aggregate reading is this model-specific, does the *decision* a screen would make move with it? We treat the dark interface as one *design* and ask whether the six backends agree on its risk (Figure 2). They do not. Aggregate wrong-advice adoption for the identical interface is 0.20–0.55 across the six beer backends and 0.12–0.62 across the amzbook backends—a spread of 0.35 and 0.50 on a 0–1 scale. This range is the threshold-free heart of the non-invariance: whatever operating threshold τ a team picks in the plausible band (0.35–0.55 beer, 0.30–0.62 amzbook), some backends land above it and others below, so no threshold in that band makes the six backends agree. Turning the estimate into the binary decision a screen would make—“flag for a human study if adoption $\geq \tau$ ”—therefore inherits that instability. At $\tau = 0.5$ one backend (gpt-4.1) flags each interface and the other five clear it: a **pairwise backend**

³Re-estimating the pooled OR with few-cluster-valid methods (only eight matched-wrong item clusters): item random-intercept (beer 2.02, amzbook 1.22, CI 0.93–1.61), item fixed effects (beer 1.92), and an item-cluster bootstrap (beer 1.93). Beer excludes 1 under all three; amzbook crosses 1 under the random-intercept model.




Fig. 3. Backend disagreement is not an artifact of the $\tau = 0.5$ operating point. Sweeping the decision threshold, the pairwise backend flip rate stays positive across the whole plausible band on both datasets (nonzero for $\tau \in [0.20, 0.55]$ beer and $[0.13, 0.62]$ amzbook), peaking near $\tau \approx 0.43$. The $\tau = 0.5$ we plot elsewhere (dotted line) is one operating point on a wide plateau, not a knife-edge.

flip rate—the fraction of backend pairs that disagree on the flag—of 0.33 on both datasets. But this point count is itself fragile: gpt-4.1’s Wilson interval reaches below 0.5 and several clearing backends’ intervals reach above it (Figure 2), which is exactly why we rest the non-invariance claim on the *range* rather than on a lucky threshold. Sweeping τ across its full range makes this explicit (Figure 3): the pairwise flip rate stays positive across the entire plausible band—and well beyond it—so the disagreement is a property of the backends, not an artifact of the single $\tau = 0.5$ operating point we plot in Figure 2. We therefore read any single- τ flip rate as *decision instability conditional on whatever operating threshold a team adopts*—not as a calibrated safety classifier, since τ itself is not validated against human harm. Two threshold-free companions tell the same story without fixing τ : the raw adoption range 0.12–0.62 above, and the persona-level risk flags, which agree only moderately across backends (Fleiss’ $\kappa = 0.55$ beer, 0.45 amzbook). In measurement terms, the instrument is not invariant: swapping the backend can move an interface from “high risk” to “low risk” without touching the design. A behavior audit rules out one benign explanation—that low-adopting backends simply *refuse* the deceptive task: explicit challenge/refusal language is rare (0% of dark responses for the OpenAI backends, 2–4% for claude and gemini), and low adoption is achieved by evidence-based re-verification, not by declining to answer.

5.4 Everything the analyst chooses moves the number

[**Mechanism (RQ2).**] So who is really moving the number—the model, or the researcher? A variance decomposition gives an uncomfortable answer: *both*, and the researcher more. A crossed random-effects model of per-trial adoption

(adopt $\sim 1 + (1|\text{backend}) + (1|\text{persona}) + (1|\text{item})$) apportion variance *more* to the researcher-chosen items and personas than to the backend (approximate variance shares: item ≈ 0.28 , persona ≈ 0.23 , backend ≈ 0.07 , residual ≈ 0.42 ; between-group SDs 0.16/0.25/0.13 corroborate the ordering)—which sharpens rather than contradicts the decision flips above: the backend is a modest source of *per-trial* noise but still shifts the *aggregate* estimate a practitioner reads enough to flip the screen. The practical implication is the same for every source: item selection, persona wording, and backend are all analyst choices that move the reading, so researchers should report a synthetic-panel result as a *distribution* over these choices rather than a single point estimate. Generation randomness—the one source a single draw per prompt cannot expose—turns out to be the smallest of all. In a preregistered multi-generation run (three backends spanning the adoption range $\times$ two UI conditions $\times$ six personas $\times$ 20 items $\times$ five generation seeds; 3,600 trials), generation contributes essentially none of the variance in dark adoption (fixed-effects $\eta^2 = 0.001$, versus 0.21 for items, 0.15 for personas, and 0.07 for the backend; Figure 4; a separate decomposition of this smaller run). Each backend's aggregate over-reliance estimate moves by at most SD = 0.03 across the five seeds (flip-metric range ≤ 0.09), small against the 0.31 spread *between* backends that flips the decision. Generation noise is real at the level of an individual persona-item cell (mean cell SD 0.17; about a third of cells change across seeds), but it averages out: two-thirds of cells are reproduced identically. The aggregate reading a practitioner acts on is therefore stable across generations, and the backend and persona effects above are not artifacts of a lucky draw.⁴

5.5 A backend-stable manipulation check: the panel propagates a prompted trust disposition

[**Mechanism (RQ3): manipulation check.**] Is anything about the panel stable? One thing is. The aggregate adoption rate is model-idiosyncratic (range 0.35 across beer backends, 0.50 across amzbook backends), but the *ordering* of which persona adopts most holds steady (Figure 5): an **explicitly trusting-novice** persona (p5) is the strict top adopter in **all 12** backend $\times$ dataset cells. Because items are reused across personas and the 12 cells share prompts, items, and providers, we do not treat trials as independent; a crossed random-effects logistic model with item, backend, and dataset random effects, together with an item-clustered GEE, gives a large odds ratio (≈ 17 , 95% CI 14–20) for that persona relative to the others—though with p5 near the ceiling in most cells, the model-free ranking rather than the OR *magnitude* is the evidence. We read this as a manipulation check the panel *passes*, not as a headline discovery: p5 is *prompt-defined* to trust, so “it adopts the AI most” demonstrates faithful, stable propagation of a prompted disposition—a reliability property of the simulation, not a discovered human vulnerability. Beneath that shared peak the vendors agree less: mean pairwise Spearman across independent vendors falls to 0.73 on beer and 0.38 on amzbook once p5 is excluded (0.85 and 0.66 with it), and the ordering detail, a background-only descriptor ablation, and a non-policy dispositional control—which shows the effect survives with no trust instruction given at all—are in Appendix B.

5.6 Capability-vulnerability mismatch: strong backends under-represent the at-risk user

[**Mechanism: capability-vulnerability.**] Now the study's most unsettling turn. A backend's aggregate adoption is one thing; *which* users it flags as at-risk is another—and combining the two exposes an uncomfortable pattern (Figure 6). Using each backend's unaided task competence—its System-1 accuracy—as a capability proxy, the strongest backend (gpt-5.5, 0.94) is precisely the one that least reproduces the highest-severity failure: on the trusting-novice persona (p5), dark-condition adoption is ≈ 1.0 for five backends but only **0.42** for gpt-5.5 (0.54 beer, 0.30 amzbook). Isolating AI influence from baseline error—the *AI-induced flip*, $P(\text{adopt} \mid \text{System-1 correct})$, which excludes trials where the

⁴We report a fixed-effects η^2 share because a random-effects variance component for the backend is degenerate on three groups; the direct stability statistics are the primary, model-free evidence.




Fig. 4. What moves a synthetic-panel risk number. Variance shares (η^2) from the preregistered multi-generation run (three backends $\times$ two UI conditions $\times$ six personas $\times$ 20 items $\times$ five generation seeds; 3,600 trials). The analyst-chosen item and persona dominate; the backend is a modest share; the random generation draw is essentially zero ($\eta^2 = 0.001$). A modest per-trial backend share still shifts the aggregate estimate enough to flip the decision (Section 5.3).

Dark-condition adoption by persona and backend (personas sorted by susceptibility; p5 highlighted)




Fig. 5. Persona $\times$ model wrong-advice adoption on the dark condition, both datasets. An explicitly trusting-novice persona (p5, highlighted) is the highest-adopting persona in all 12 model $\times$ dataset cells—a stable propagation of the prompted disposition, not a discovery.

persona was already wrong (there adoption is 1 by construction)—the same ordering holds: the frontier flips a correct answer to the wrong AI least often (0.11), a mid-tier backend most (0.44). In short, *you would not crash-test with your*

Manuscript submitted to ACM

The frontier backend is nearly blind to the at-risk persona (coverage, not the fit, carries this)




Fig. 6. **The frontier backend under-represents the at-risk user.** *Left:* backend task competence (System-1 accuracy) vs. the AI-induced flip rate $P(\text{adopt} | \text{System-1 correct})$, both datasets—a directional negative trend (Spearman $-0.26 / -0.54$; n.s. at $n = 6$, so the fit is a visual guide, not the load-bearing evidence). *Right:* vulnerability coverage (number of the six personas whose dark adoption reaches 0.5) for a single frontier backend, a weak pair, and the full panel—the model-free reading the argument rests on: the single frontier backend surfaces the fewest at-risk profiles.

safest driver—yet a synthetic panel powered by the frontier model does exactly that, screening for over-reliance danger with the one “user” least capable of being over-reliant.

The pattern sharpens at $n = 11$ backends. It is not an artifact of the six-model ladder. Widening to **eleven** backends across three vendors and a broad capability range, and pooling both datasets, capability anti-correlates with all three vulnerability measures (Spearman: aggregate adoption -0.82 , AI-induced flip -0.67 , p5 adoption -0.65 ; Benjamini–Hochberg $p \leq 0.03$; Figure 7).⁵ The aggregate-adoption anti-correlation is the sturdy one—significant in each dataset ($-0.68 / -0.98$) and every leave-one-backend-out jackknife; the flip and p5 correlations are amzbook-carried and single-backend fragile. But the reading does not hang on the correlation: **the frontier backend is the least sensitive to the at-risk persona.**

This has a concrete consequence for assembling an evaluation panel. Treating a persona as *covered* when its dark adoption reaches a risk threshold (0.5), a single frontier backend—the model a quality-optimal router [19] would select—surfaces the risk for only 1 of 6 personas on beer and 0 of 6 on amzbook, whereas a pair of weaker backends (gpt-4.1, gpt-4o-mini) surfaces 4 and 3 (matching the full six-backend panel). Coverage is monotone in panel size by construction, so the point is not “more models is better”; it is that the *default best choice—the strongest model—provides the least coverage of the failure cells*. Choosing an evaluation panel to maximize **synthetic-vulnerability coverage** is therefore the *inverse* of choosing it to maximize task quality: a practice built on “use the strongest model” systematically misses the failure cells that matter most.

⁵Added backends: gpt-3.5-turbo, gpt-4, gpt-5.4, claude-opus-4.6, and gemini-3.1-pro; a twelfth, claude-haiku-4.5, is excluded because it never emitted a parseable decision—an instrument failure, disclosed.

Backend expansion ($n = 11$): capability anti-correlates with vulnerability; coverage routing still inverted




Fig. 7. Backend expansion to $n = 11$ (pooling both datasets). Left: backend task competence (System-1 accuracy) vs. three dark-condition vulnerability measures, tiered by robustness—aggregate adoption (filled, solid fit) is jackknife-robust and significant in each dataset; AI-induced flip and p5 adoption (open, dashed fit) are BH-significant only pooled and are amzbook-carried / single-backend fragile. Right: the $n = 11$ vulnerability-coverage curve—a capability-first (router) panel reaches full coverage of the high-risk cells only at $k = 10$, while a coverage-greedy panel is complete almost immediately; a single frontier backend covers 1/9.

We can make this selection principle precise. Let C be the set of high-severity cells (dataset $\times$ persona; high-severity by the *panel's own* adoption) and $a_b(c)$ backend b 's dark adoption on cell c ; a panel S of backends *covers* every cell that any of its members reproduces above a risk threshold τ , with coverage

$$f(S) = |\{c \in C : \max_{b \in S} a_b(c) \geq \tau\}|. \quad (2)$$

The *coverage-greedy* panel adds, at each step, the backend that covers the most still-uncovered high-severity cells.⁶ Figure 7 (right) shows this contrast at $n=11$; the same $n=6$ comparison of a *capability-first* panel that adds backends strongest-first, i.e. the order a quality-optimal router [19] would prefer, is almost opposite to the coverage-greedy panel. On the clean AI-induced-flip metric the greedy panel covers every high-risk cell with two backends, whereas the capability-first panel's first pick—the frontier model—covers a single cell and only reaches full coverage at $k = 6$. Under the raw-adoption metric the contrast is starker still: one *low-capability* backend (gpt-4.1) already covers all seven high-risk cells, so the single best probe is the weakest model and the router's frontier pick covers the fewest cells. Selecting a panel for vulnerability coverage is thus not merely different from capability routing—it is close to its inverse.

A caveat is built into this metric, and we state it plainly: a cell counts as high-severity by the panel's own adoption, and a backend counts as covering it by that same adoption, so higher-adopting backends cover more *by construction*. Coverage is a *sensitivity* construct, not a validity one—at the limit, a backend that adopts everything has perfect coverage and no discriminative value. What is *not* definitional is *why* the weaker models adopt more: capability and adoption anti-correlate empirically, so the router's default—the strongest model—is the *least sensitive* prober.

⁶ f is monotone and submodular, so this greedy order carries the standard $(1 - 1/e)$ approximation guarantee against the best size- k panel.




Fig. 8. Coverage buys sensitivity with specificity. For each backend, wrong-advice adoption on the coercive (dark) interface—the sensitivity that coverage rewards—plotted against wrong-advice adoption on the *non*-coercive (neutral) interface, a false-flag rate on a design that carries no dark pattern. The two track each other tightly (Spearman 0.86): the same high-coverage backends (gpt-4.1) also over-adopt on safe interfaces, while the frontier backend (gpt-5.5) is low on both. Points sit above the diagonal, so coercion does add signal beyond baseline credulity—but a backend that “flags everything” would sit *on* it, which is why coverage alone cannot certify an instrument.

This trade-off is measurable, and it guards against the obvious failure mode: a backend that simply adopts every wrong AI would post perfect coverage. It does not come free. The backends that cover the most dark cells are also the ones that adopt the most wrong advice on the *non*-coercive (neutral) interface—a false alarm, since that design carries no dark pattern—with a Spearman of 0.86 between the two (Figure 8). Coverage therefore buys sensitivity at the cost of specificity, exactly as a screening trade-off should; the coercion-specific signal (dark minus neutral adoption)

A backend-aware protocol for synthetic-user interface screening. Before treating a synthetic panel’s risk reading as a property of the design:

- (1) **Power the panel for coverage, not capability.** Use several capability-diverse backends chosen to surface the failure modes a single strong model misses—not the one top-ranked model (Section 5.6).
- (2) **Report a distribution, not a point.** Give each backend’s estimate with an interval; never collapse to one number (Section 5.4).
- (3) **Measure decision instability.** Compute the pairwise decision-flip rate across backends at your operating threshold, and report the threshold-free adoption range alongside it (Section 5.3).
- (4) **Ablate the persona policy.** Re-run with the explicit trust/deference instruction removed, to separate a prompt-compliance artifact from a dispositional signal (Appendix B).
- (5) **Abstain on disagreement.** When backends disagree beyond a preset margin (illustratively, range > 0.2), decline to certify and route the design to a human study.
- (6) **Prefer relative to absolute.** For design comparisons, report the safeguard’s effect *across* backends and treat a single-backend “it helps”/“it doesn’t” as provisional (Appendix E).

Fig. 9. The reporting protocol our audit motivates, as a step-by-step screening workflow. It presumes no human ground truth; step 6 is where a human study (Section 8) takes over.

stays positive for every backend, so the dark pattern does add signal beyond baseline credulity. Read correctly, then, the selection principle is a *conservative* screening-design choice—favour sensitivity when a missed at-risk cell costs more than a false alarm—not a claim that credulous backends are simply better; whether higher sensitivity improves screening against real users is what the planned study (Section 8) must decide.

The capability–vulnerability inversion is also not merely mechanical. A more-accurate backend cannot adopt a wrong AI on a trial it already answers correctly, so higher System-1 accuracy mechanically depresses raw adoption. To remove that advantage we restrict to the cells where *all six* backends’ System-1 answer is correct (beer $n=133$, amzbook $n=131$ item×persona cells) and recompute adoption there: the frontier backend still adopts far less (gpt-5.5 0.06–0.08) than the weaker models (0.24–0.42), so a robust residual inversion survives accuracy matching. Conditioning the AI-induced flip on a correct prior tells the same story (gpt-5.5 0.14/0.08 vs. gpt-4.1 0.43/0.46). The mechanical component explains part of the raw gap, not the direction: even head-to-head on identical, already-correct trials, the strongest backend is the least likely to be talked out of a right answer by a coercive interface.

The same mismatch appears in the vocabulary of appropriate reliance [25]: in the field’s own terms the frontier backend is the *most appropriately reliant* user—it has the highest relative self-reliance, keeping its own correct answer when the AI is wrong, and the lowest over-reliance under the adversarial dark condition—which is precisely what makes it the least useful probe for an interface whose danger *is* over-reliance. Capability trends with appropriate reliance across the six backends but is not significant at $n = 6$, so we read it as directional support for what the coverage result shows directly (Appendix C).

The ordering, and the inversion that produces it, also replicate on a structurally different LSAT logical-reasoning task—4-way multiple choice, harder, and a different advice format (Appendix D).

6 Discussion

The backend is part of the measurement instrument. A synthetic panel does not simply reveal an interface’s risk; it co-produces that reading with the backend, prompt, and generation the analyst chooses. Changing the backend therefore changes the instrument. When that choice flips the interface’s risk classification (Section 5.3), a single-backend

“design risk” number is not yet a property of the design. The question “which LLM should power my synthetic study?” is not an implementation detail; it is a first-order methodological decision.

Capability is not validity. Choosing the strongest model as the best simulator is tempting, but our frontier model follows a coercively-framed wrong AI least. That may make it a better *task-solver*; it does not establish that it better tracks the target users—you would not crash-test with your safest driver. Higher capability can even compress the persona differentiation on which the screen relies. Without a human criterion, capability and correctness cannot justify the choice of synthetic participant. Nor do our results show that the frontier model “misses real danger” or that weaker models are “more human”—either claim requires the human study we have not run.

Disagreement should trigger abstention, not averaging. Cross-backend disagreement is not merely measurement error. It identifies the interfaces for which a single-backend verdict is least trustworthy. Treating that disagreement as information turns our reporting protocol from general caution into an actionable rule: when the instruments disagree, the panel should decline to certify (Figure 9 collects these steps into a single workflow).

Three kinds of failure. Our evidence separates (i) an *interface-level synthetic signal* (agents show adoption or disagreement), from (ii) a *measurement failure* (the signal flips when the backend or prompt changes), from (iii) a *human-facing design failure* (real users are harmed). This paper establishes (i) and characterizes (ii); (iii) is reserved for the human study.

Implications for commercial synthetic-user tooling. Our findings speak directly to two assumptions that a synthetic-user platform following standard engineering defaults would plausibly adopt [30, 32]. *First, rethink “frontier-first.”* The instinct to power a screen with the strongest available model—the choice a quality-optimal router makes [19]—is close to the least sensitive choice for a risk audit: the strongest backend least reproduces the highest-severity failure (Section 5.6), so a risk panel should be assembled for *synthetic-vulnerability* coverage, which can require deliberately including weaker or more diverse backends: a *single* frontier backend covers almost none of the risk that a weak pair recovers (Section 5.6). *Second, retire the single point estimate.* A one-backend number (“wrong-advice adoption = 0.6”) is not backend-invariant: it moves enough to flip the screening decision (Section 5.3). Like a weather forecast, researchers should report a synthetic-panel result as a *distribution* over backends, prompts, and generations, with a disagreement-triggered *abstention*: when backends disagree beyond a preset margin (here, illustratively and not yet calibrated, a range exceeding 0.2), the tool should decline to certify and route the design to human testing rather than emit a safe/unsafe verdict. *Third, prefer relative to absolute claims.* Even for the comparative question such tools are usually asked—“does this safeguard help?”—the *direction* of the improvement is backend-robust while its magnitude, and on one backend its very presence, is not, so a single-backend “it helps”/“it doesn’t” is provisional (Appendix E). None of this replaces a human study; it makes an honest synthetic screen a *triage* step that knows when to defer to one.

7 Limitations and Threats to Validity

- **Synthetic personas are not people**—the central threat. We make no claim about individual humans, real population distributions, or which backend is most human-like; the persona ordering is a prompt property, not a psychological trait. Only the human study (Section 8) grounds these.
- **The measurement is not backend-invariant.** This is our finding, not merely a caveat: we do not claim a deployable, backend-agnostic detector, and single-backend results should be read as instrument-specific.

- **Correspondence is unvalidated and underpowered** ($n = 5$ conditions; amzbook degenerate). No predictive-validity claim is made.
- **Two review datasets, both binary sentiment.** Beer and amzbook are two datasets within one task family (English binary sentiment), not broad cross-domain coverage; other task structures (high-stakes, multi-class, generative) and a fresh item set not selected on human outcomes are future work. A preregistered LSAT arm (Appendix D; 4-way multiple choice) partially closes this: the capability–vulnerability mismatch and the persona vulnerability ordering ($\rho = 0.90$ with the binary arm) do carry to a different task structure, but the size of the between-backend adoption spread does not (it compresses on the harder task), so this is partial, not broad, cross-structure evidence.
- **Coercion is a bundled manipulation.** The dark condition varies confidence, expert authority, and accountability cues together, so the framing effect (Section 5.2) is a *bundled* coercive effect; we do not factorially separate its components.
- **Generation and prompt-format sensitivity.** The main results use one generation per prompt; we estimate generation stochasticity directly in a preregistered multi-generation run (Section 5.4) and find it negligible for the aggregate estimate ($\eta^2 = 0.001$; per-backend SD ≤ 0.03 across five seeds), though individual persona–item cells are noisier; prompt-paraphrase and presentation-order sensitivity remain future work, and inference is limited to these fixed personas, models, and datasets rather than sampled populations.
- **Partial specification audit.** We decompose backend/persona/item variance and sweep the decision threshold, but a full specification-curve/multiverse analysis [29] over parsing rules, heterogeneity metrics, and conflict-conditioning choices is future work.
- **Item selection uses human reliance variance** (disclosed). Holding items out from threshold fitting does not remove selection on the human criterion; a fresh-item human check remains, so we treat this as unresolved pending the human study.
- **Protective interventions are textual framings, underpowered for backend contrasts.** The cognitive-forcing and verification conditions (Appendix E) are prompt-level renderings of the interventions, on one domain and one generation; the backend $\times$ condition interaction is underpowered, so we claim sign-robustness of the protective effect (all benefits ≥ 0), not magnitude-invariance—and note one backend (gemini) where the safeguard is nearly inert.
- **The capability–vulnerability mismatch rests on synthetic evidence.** Capability is a coarse proxy (mean System-1 accuracy) and part of the raw capability $\times$ adoption correlation is mechanical (a backend that answers correctly cannot then adopt a wrong AI); the finer flip/p5 correlations are amzbook-carried and single-backend fragile, which is why the load-bearing result is the assumption-light coverage inversion rather than the correlation. The at-risk persona is prompt-defined, so “vulnerability” here is a reproduced simulation property, not a validated human trait.
- **Prompt-level, not pixel-level.** Agents see a textual encoding of the interface, so we study prompt framing rather than visual interface effects; a multimodal rendering is future work.

8 Planned Human Validation

This study is designed but not yet run or preregistered. The one thing synthetic compute cannot buy is grounding against real people. We have designed a stripped, axis-2-only human study (within-subjects, $N \approx 40$, Prolific) that would test, against humans and across the two review datasets, whether the coercive framing increases wrong-advice adoption

relative to a faithful/placebo baseline, whether dispositional deference (high self-reported AI-deference $\times$ low domain skill) moderates that effect, and which backend or aggregation rule best approximates the human pattern. We make no prediction-validity claim here—without a human criterion we cannot say which backend is the most human-faithful screener; all outcomes, including nulls, would be publishable bounds on synthetic-panel validity.

9 Conclusion

An AI-advice interface can fail by coercing over-reliance on a confidently wrong AI—and, more diffusely, by making reliance dangerously uneven across users. We audit the first. Synthetic panels promise to catch such failures before an expensive human study. Our audit shows why that promise must be qualified: the panel's risk reading is **backend-dependent**.

For the *same* interface, aggregate wrong-advice adoption swings widely across backends—enough to flip the screening decision across most plausible thresholds. Yet a matched contrast still identifies a real coercive-framing effect whose backend-dependence we cannot detect. The backend moves the adoption *level*; any effect on the framing slope is below what our data resolve. Backend choice is therefore part of the methodology, with direct consequences for screening.

Beneath this instability, the panel reliably propagates one prompted disposition: an explicitly trusting-novice persona ranks highest across every backend and dataset. That consistency is a property of the simulation, not a discovery of who is at risk. Nor is the strongest model automatically the most valid simulator: the frontier backend is least sensitive to the panel's highest-severity failure. Synthetic screens should therefore report distributions across backends, prompts, and generations, and abstain when their instruments disagree.

We offer an honest map of where synthetic-user screening can and cannot be trusted today. Until the human study is run, the most trustworthy synthetic panel is not one that always answers. It is one that knows when to abstain.

References

- [1] Gati Aher, Rosa I. Arriaga, and Adam Tauman Kalai. 2023. Using Large Language Models to Simulate Multiple Humans and Replicate Human Subject Studies. In *Proceedings of the 40th International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 202)*. PMLR, 337–371. <https://proceedings.mlr.press/v202/aher23a.html>
- [2] Lisa P. Argyle, Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. 2023. Out of One, Many: Using Language Models to Simulate Human Samples. *Political Analysis* 31, 3 (2023), 337–351. doi:10.1017/pan.2023.2
- [3] Gagan Bansal, Tongshuang Wu, Joyce Zhou, Raymond Fok, Besmira Nushi, Ece Kamar, Marco Tulio Ribeiro, and Daniel S. Weld. 2021. Does the Whole Exceed its Parts? The Effect of AI Explanations on Complementary Team Performance. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems (CHI '21)*. ACM, 1–16. doi:10.1145/3411764.3445717
- [4] James Bisbee, Joshua D. Clinton, Cassy Dorff, Brenton Kenkel, and Jennifer M. Larson. 2024. Synthetic Replacements for Human Survey Data? The Perils of Large Language Models. *Political Analysis* 32, 4 (2024), 401–416. doi:10.1017/pan.2024.5
- [5] Jessica Y. Bo, Sophia Wan, and Ashton Anderson. 2025. To Rely or Not to Rely? Evaluating Interventions for Appropriate Reliance on Large Language Models. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. ACM, 1–23. doi:10.1145/3706598.3714097
- [6] Zana Bućinca, Maja Barbara Malaya, and Krzysztof Z. Gajos. 2021. To Trust or to Think: Cognitive Forcing Functions Can Reduce Overreliance on AI in AI-Assisted Decision-Making. *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW1 (2021), 1–21. doi:10.1145/3449287
- [7] Yong Cao, Haijiang Liu, Arnav Arora, Isabelle Augenstein, Paul Röttger, and Daniel Hershcovich. 2025. Specializing Large Language Models to Simulate Survey Response Distributions for Global Populations. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. 3141–3154. doi:10.18653/v1/2025.naacl-long.162
- [8] Myra Cheng, Tiziano Piccardi, and Diyi Yang. 2023. CoMPoS: Characterizing and Evaluating Caricature in LLM Simulations. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP '23)*. 10853–10875. doi:10.18653/v1/2023.emnlp-main.669
- [9] Ricardo Dominguez-Olmedo, Moritz Hardt, and Celestine Mender-Dünner. 2024. Questioning the Survey Responses of Large Language Models. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2306.07951.
- [10] Krzysztof Z. Gajos and Lena Mamykina. 2022. Do People Engage Cognitively with AI? Impact of AI Assistance on Incidental Learning. In *Proceedings of the 27th International Conference on Intelligent User Interfaces (IUI '22)*. ACM, 794–806. doi:10.1145/3490099.3511138

- [11] Jennifer Haase, Jana Gonnemann-Müller, Paul H. P. Hanel, Nicolas Leins, Thomas Kosch, Jan Mendling, and Sebastian Pokutta. 2026. Within-Model vs Between-Prompt Variability in Large Language Models for Creative Tasks. arXiv:2601.21339. arXiv preprint.
- [12] Perttu Härmäläinen, Mikke Tavast, and Anton Kunnari. 2023. Evaluating Large Language Models in Generating Synthetic HCI Research Data: A Case Study. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems (CHI '23)*. ACM, 1–19. doi:10.1145/3544548.3580688
- [13] Tiancheng Hu and Nigel Collier. 2024. Quantifying the Persona Effect in LLM Simulations. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers) (ACL '24)*. 10289–10307. doi:10.18653/v1/2024.acl-long.554
- [14] Vivian Lai and Chenhao Tan. 2019. On Human Predictions with Explanations and Predictions of Machine Learning Models: A Case Study on Deception Detection. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* '19)*. ACM, 29–38. doi:10.1145/3287560.3287590
- [15] Jamie Luguri and Lior Jacob Strahilevitz. 2021. Shining a Light on Dark Patterns. *Journal of Legal Analysis* 13, 1 (2021), 43–109. doi:10.1093/jla/laaa006
- [16] Thomas Mildner, Gian-Luca Savino, Philip R. Doyle, Benjamin R. Cowan, and Rainer Malaka. 2023. About Engaging and Governing Strategies: A Thematic Analysis of Dark Patterns in Social Networking Services. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems (CHI '23)*. ACM, 1–15. doi:10.1145/3544548.3580695
- [17] Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timmit Gebbru. 2019. Model Cards for Model Reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT*)*. ACM, 220–229. doi:10.1145/3287560.3287596
- [18] Moran Mizrahi, Guy Kaplan, Dan Malkin, Rotem Dror, Dafna Shahaf, and Gabriel Stanovsky. 2024. State of What Art? A Call for Multi-Prompt LLM Evaluation. *Transactions of the Association for Computational Linguistics* 12 (2024), 933–949. doi:10.1162/tacl_a_00681
- [19] Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez, M. Waleed Kadous, and Ion Stoica. 2025. RouteLLM: Learning to Route LLMs with Preference Data. In *Proceedings of the 13th International Conference on Learning Representations (ICLR)*. arXiv:2406.18665.
- [20] Joon Sung Park, Lindsay Popowski, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. 2022. Social Simulacra: Creating Populated Prototypes for Social Computing Systems. In *Proceedings of the 35th Annual ACM Symposium on User Interface Software and Technology (UIST '22)*. ACM, 1–18. doi:10.1145/3526113.3545616
- [21] Joon Sung Park, Carolyn Q. Zou, Jonne Kamphorst, Niles Egan, Aaron Shaw, Benjamin Mako Hill, Carrie Cai, Meredith Ringel Morris, Percy Liang, Robb Willer, and Michael S. Bernstein. 2024. LLM Agents Grounded in Self-Reports Enable General-Purpose Simulation of Individuals. arXiv:2411.10109v3. arXiv preprint; revised June 2026.
- [22] Charvi Rastogi, Yunfeng Zhang, Dennis Wei, Kush R. Varshney, Amit Dhurandhar, and Richard Tomsett. 2022. Deciding Fast and Slow: The Role of Cognitive Biases in AI-Assisted Decision-Making. *Proceedings of the ACM on Human-Computer Interaction* 6, CSCW1 (2022), 1–22. doi:10.1145/3512930
- [23] Sasha Robinson, Katherine M. Collins, Ilia Sucholutsky, and Kelsey R. Allen. 2026. Under the Influence: Quantifying Persuasion and Vigilance in Large Language Models. arXiv:2602.21262v3. arXiv preprint; revised March 2026.
- [24] Shibani Santurkar, Esin Durmus, Faisal Ladhak, Cinoo Lee, Percy Liang, and Tatsunori Hashimoto. 2023. Whose Opinions Do Language Models Reflect?. In *Proceedings of the 40th International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 202)*. PMLR, 29971–30004. <https://proceedings.mlr.press/v202/santurkar23a.html>
- [25] Max Schemmer, Niklas Kuehl, Carina Benz, Andrea Bartos, and Gerhard Satzger. 2023. Appropriate Reliance on AI Advice: Conceptualization and the Effect of Explanations. In *Proceedings of the 28th International Conference on Intelligent User Interfaces (IUI '23)*. ACM, 410–422. doi:10.1145/3581641.3584066
- [26] Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. 2024. Quantifying Language Models' Sensitivity to Spurious Features in Prompt Design, or: How I Learned to Start Worrying about Prompt Formatting. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*. arXiv:2310.11324.
- [27] Preethi Seshadri, Samuel Cahyawijaya, Ayomide Odumakinde, Sameer Singh, and Seraphina Goldfarb-Tarrant. 2026. Lost in Simulation: LLM-Simulated Users are Unreliable Proxies for Human Users in Agentic Evaluations. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 47423–47439. doi:10.18653/v1/2026.acl-long.2192
- [28] Joseph P. Simmons, Leif D. Nelson, and Uri Simonsohn. 2011. False-Positive Psychology: Undisclosed Flexibility in Data Collection and Analysis Allows Presenting Anything as Significant. *Psychological Science* 22, 11 (2011), 1359–1366. doi:10.1177/0956797611417632
- [29] Uri Simonsohn, Joseph P. Simmons, and Leif D. Nelson. 2020. Specification Curve Analysis. *Nature Human Behaviour* 4 (2020), 1208–1214. doi:10.1038/s41562-020-0912-z
- [30] Synthetic Users. 2026. Synthetic Users, Organic Insights. <https://www.syntheticusers.com/>. Commercial synthetic-user research platform; accessed July 29, 2026.
- [31] Lindia Tjautja, Valerie Chen, Tongshuang Wu, Ameet Talwalkar, and Graham Neubig. 2024. Do LLMs Exhibit Human-like Response Biases? A Case Study in Survey Design. *Transactions of the Association for Computational Linguistics* 12 (2024), 1011–1026. doi:10.1162/tacl_a_00685
- [32] Uxia. 2026. 7 Top Synthetic User Testing Platforms to Watch in 2026. <https://www.uxia.app/blog/7-top-synthetic-user-testing-platforms-to-watch-in-2026>. Industry survey of commercial synthetic-user (AI persona) UX-testing tools; accessed July 29, 2026.
- [33] Helena Vasconcelos, Matthew Jörke, Madeleine Grunde-McLaughlin, Tobias Gerstenberg, Michael S. Bernstein, and Ranjay Krishna. 2023. Explanations Can Reduce Overreliance on AI Systems During Decision-Making. *Proceedings of the ACM on Human-Computer Interaction* 7, CSCW1 (2023), 1–38. doi:10.1145/3579605
- [34] Yunfeng Zhang, Q. Vera Liao, and Rachel K. E. Bellamy. 2020. Effect of Confidence and Explanation on Accuracy and Trust Calibration in AI-Assisted Decision Making. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency (FAT* '20)*. ACM, 295–305.

doi:10.1145/3351095.3372852

- [35] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*. arXiv:2306.05685.

A Supplement: The Over-Dispersion Axis (Candidate Diagnostic)

The second axis of our construct—*between-user reliance over-dispersion* (Section 3)—is reported here rather than in the main paper, because its human counterpart is null in every configuration we can test. We treat it as a candidate diagnostic, not a load-bearing result.

The human failure is real, and measurable. We first verify that over-dispersion corresponds to a measurable human quantity in the source benchmark. On the Bansal et al. [3] human data, we estimate between-user reliance over-dispersion with the same beta-binomial estimator we apply to the panel; conflict-conditioned over-dispersion is positive and excludes zero ($\rho \approx 0.067$). By split-half reliability—the primary well-identified estimator on this single-observation-per-cell structure—no-AI reliance behaves as a stable trait (stable-user share ≈ 0.74), while under AI assistance that share falls to a task-dominated 0.32–0.41. (An optional variance-components GLMM did not converge and is not used; one AI condition, expert-adaptive, yields a degenerate, negative split-half reliability and is excluded from that range.) We report this anchor briefly (full participant counts, per-condition estimates, and intervals are in the supplement); it establishes that the human construct is measurable but does *not* generalize to a second human dataset (Lu & Yin), which we disclose as a boundary. The human anchor is not used to validate the synthetic panel—that link is null in every configuration we test (below).

Synthetic persona disagreement is backend-dependent. The panel's internal disagreement is similarly backend-dependent (Figure 10). We call this quantity *synthetic persona disagreement*—the between-persona variance of the six personas' conflict-conditioned reliance rates (Section 3)—and are explicit that it measures how differently our *prompted* personas respond, not a validated human over-dispersion. Mean persona disagreement across the model set is 0.11 / 0.16 / 0.10 / **0.012** (beer) and 0.12 / 0.19 / 0.11 / **0.015** (amzbook): gpt-5.5 is the most homogeneous on both datasets, and the beta-binomial ρ (which mean disagreement tracks) is correspondingly near zero. Disagreement is computed over the full 5-condition design (5 conditions $\times$ 6 personas $\times$ 20 items = 600 trials/model), not the dark condition alone; the low value at gpt-5.5 is not a no-conflict artifact, since genuine conflict trials remain present (e.g. 215/600 on amzbook), and it also holds on the common conflict subset shared across models. We caution against reading it as a safety construct collapsing: low disagreement at a strong model can equally reflect the personas converging on the *correct* answer (uniform task-solving) rather than an interface failing to differentiate users—these are indistinguishable without a human criterion.

Panel $\leftrightarrow$ human correspondence is unvalidated (honest null). Across every configuration we can test, the panel's per-condition axis-1 signal does *not* significantly track the human anchor's (Spearman on $n = 5$ conditions; null under BH for all six models). On amzbook the correspondence is *degenerate*—the human anchor's per-condition over-dispersion is near-constant, so the test is uninformative rather than a genuine null. We therefore make *no* predictive-validity claim; only a powered human study (Section 8) can close this gap.




Fig. 10. Axis-1 mean panel disagreement across the same-provider model set, both datasets (error bars: SD across the 5 UI conditions). gpt-5.5 shows the lowest cross-persona disagreement on both datasets.

B Supplement: Manipulation-Check Robustness

The main paper reports the manipulation check itself (Section 5.5): an explicitly trusting-novice persona (p5) is the strict top adopter in all 12 backend×dataset cells. This section carries the checks that locate that signal and the cross-vendor ordering detail beneath it.

Is the effect an artifact of the “trust the AI” wording? Not on its own. Holding the “novice” background fixed, the novice-*skeptical* persona adopts far less than the novice-*trusting* one on every backend, so the effect tracks the disposition rather than the “novice” label. A background-only *descriptor ablation* (persona background stated but no trust/deference wording; two backends, 20 items) retains the trusting-novice’s high adoption on the mid-tier backend (1.00 with policy vs. 0.90 without) but shows it is more policy-carried on the frontier (p5 0.60 → 0.40, rank 1→2), with skeptical personas also becoming more compliant once their verification instruction is removed—so part of the effect is prompt-compliance and part is background, in a backend-dependent mix (20 items, one generation, per-cell SE ≈ 0.11; individual rank swaps read cautiously). Two alternative inductions (a first-person backstory and a few-shot behavioural demonstration) keep p5 the top pooled adopter (1.00) but do not strip the deference *content*, so we read them as re-encodings of the disposition rather than deference-free controls. Because p5 was built to be trusting, a chance baseline of 1/6 per persona would in any case be inappropriate—the personas are not exchangeable.

A non-policy dispositional control. To probe circularity more directly, we ran a further exploratory ablation with a *dispositional* persona that states three facts—low task experience, frequent AI use, and *low self-confidence in one’s*

Non-policy dispositional ablation (exploratory: 12 items, 1 generation, 4 backends)




Fig. 11. Non-policy dispositional ablation (exploratory: 12 beer items, one generation, dark condition, four capability-spanning backends). *Left*: a deferential disposition (d1: stated low task experience, frequent AI use, and *low self-confidence*, but *no* instruction to trust the AI) adopts the wrong AI more than an independent disposition (d2) on every backend; the frontier (gpt-5.4) compresses the gap. *Right*: decomposing the disposition one facet at a time around d1, low self-confidence is the consistent positive lever, while stated experience and AI-use frequency are near-inert—the panel over-relies when a persona doubts its own judgment, not merely when it is labelled a novice.

own judgment—plus a decision instruction identical across personas, but *never* an instruction to trust or defer to the AI (four capability-spanning backends: gpt-3.5-turbo, gpt-4, gpt-5.4, gemini-3.1-pro; 12 beer items; dark condition). The deferential disposition adopts the wrong AI more than an independent disposition (high experience, infrequent AI use, high self-confidence) on *all four* backends (adoption gap +0.08 to +0.83; Figure 11; compare the explicit-policy p5–p4 gap +0.50 to +0.83)—so the over-reliance signal is not merely compliance with a written “trust the AI” instruction. Decomposing the disposition one facet at a time around the deferential anchor, the lever is *low self-confidence* (positive on all four backends), while stated experience and AI-use frequency alone are near-inert: the panel over-relies when a persona doubts its own judgment, not merely when it is labelled a novice. As everywhere else, the frontier backend compresses the gap (gpt-5.4, +0.08), consistent with the capability–vulnerability mismatch (Section 5.6). This is exploratory (12 items, one generation, four backends; per-cell SE ≈ 0.14), but it strengthens the reading that the trusting-novice cell indexes a *dispositionally instantiable* failure mode, not only a prompt-defined one.

Cross-vendor ordering beneath the shared peak. The cross-vendor *ordering* agreement is largely carried by this one persona. The mean pairwise Spearman across independent vendors is 0.85 on beer (0.66 on amzbook) with p5 included, and drops to 0.73 on beer and to 0.38 on amzbook once p5 is excluded (Figure 12). Beneath the shared top rank, then, the vendors agree moderately on amzbook and still substantially on beer—the stable part of the persona profile is the peak the prompt puts there.




Fig. 12. Per-persona wrong-advice adoption profiles for the three independent-vendor backends, one panel per dataset. All vendors peak at the prompted trusting-novice persona (p5), but beneath that peak agreement weakens—removing p5, the mean pairwise Spearman falls from 0.85 to 0.73 on beer and from 0.66 to 0.38 on amzbook. Persona order is categorical (arbitrary). The shared p5 peak is a property of the persona prompt, not population discovery.

C Supplement: Appropriate Reliance (RAIR/RSR)

Section 5.6 reports the capability–vulnerability inversion in the vocabulary of appropriate reliance [25]; the definitions and numbers are here. With pre-advice and final decisions $s_1, y \in \{0, 1\}$, ground truth g , and displayed AI advice $\hat{a}$, relative positive AI reliance and relative self-reliance are the conflict-conditioned rates

$$\text{RAIR} = \Pr(y=g \mid s_1 \neq g, \hat{a}=g), \quad \text{RSR} = \Pr(y=g \mid s_1=g, \hat{a} \neq g). \quad (3)$$

On the faithful condition, where the AI is sometimes right and sometimes wrong, the frontier backend has the highest relative self-reliance (RSR = 0.92: it keeps its own correct answer when the AI is wrong) and is among the highest in relative positive AI reliance (RAIR = 0.57, second to claude’s 0.64; $n = 14$ reliance opportunities). Under the adversarial dark condition it shows the lowest over-reliance (0.16 vs. up to 0.46). Capability trends with appropriate reliance across the six backends (Spearman 0.71 with RAIR, -0.54 with dark over-reliance) but is not significant at $n = 6$, which is why the coverage inversion, not this correlation, carries the argument.

D Supplement: A Structurally Different Task (LSAT)

Could the mismatch be an artifact of one task *type*? Beer and amzbook are both English binary sentiment, so the pattern might not survive a different structure. As a preregistered robustness arm we replayed the same six backends, six personas, and dark protocol on 20 LSAT logical-reasoning items [3]—a genuinely different structure (4-way multiple choice, harder, a different advice format), where the dark AI recommends a guaranteed-wrong option and adoption is final = $\hat{a}$ against that explicit wrong option (all 720 dark trials satisfy $\hat{a} \neq g$; System-2 parse rate = 1.00 for every backend, so no instrument-failure exclusion applies). We froze a replication rule—sign of the capability×vulnerability correlation, concordance of the per-persona vulnerability *ordering* with the binary arm, and a ≥ 0.25 cross-backend spread in aggregate adoption—before computing any LSAT metric. Two of the three carry over and one does not. The persona vulnerability ordering is preserved almost exactly: the trusting-novice persona is again the top adopter (dark




Fig. 13. Protective-intervention audit. *Left*: per-backend wrong-AI adoption across dark (coercive) → plain → forcing → verify. Every backend shows a non-increasing staircase (protective benefit ≥ 0). But the magnitude ranges from ≈ 0.12 (gpt-4o-mini) to ≈ 0.01 on gemini, where the protective framing is nearly inert; the backend $\times$ condition interaction is not significant (underpowered). *Right*: the at-risk trusting-novice persona (p5, pooled) drops from 0.78 (plain) to 0.46–0.53 under protection. dark is shown descriptively (separate run).

adoption 0.54 on LSAT, 0.90 on binary) and the skilled-skeptical persona the lowest, with rank correlation $\rho = 0.90$ ($p = 0.015$, $n = 6$ personas). The capability $\times$ vulnerability relationship keeps its *sign*: the strongest backend (gpt-5.5, System-1 accuracy 0.85) again shows the lowest dark adoption (0.04) and the lowest AI-induced flip (0.05) and again flattens the trusting-novice peak (0.10 vs. 0.60–0.85 for weaker backends), and $\rho(\text{capability}, \text{flip}) = -0.83$ stays in $[-1.0, -0.7]$ under every leave-one-backend-out subsample—though its nominal $p = 0.04$ is fragile (three of six drops push it to $p = 0.19$), so we read it, as in the binary arm, as a robust *sign* rather than a significant coefficient. The coverage inversion also survives: the single highest-capability backend covers 0/6 high-adoption personas at $\tau = 0.5$ and the full panel only 1/6. What does *not* replicate is our own non-invariance magnitude: aggregate dark adoption spans only 0.04–0.27 (range 0.225, just under the frozen 0.25), because the harder 4-way task lowers wrong-advice adoption across all backends and compresses the between-backend gap. LSAT is therefore a *partial* generalization—the capability–vulnerability mismatch and the persona vulnerability ordering travel to a new task structure, while the size of the aggregate backend spread is task-dependent—not a task-invariant law.

E Supplement: Auditing a Protective Design (Relative Comparison)

Does the panel's instability doom it for practice? Not entirely: the most common use of a synthetic screen is not to certify an absolute risk number but to *compare* designs—does adding a safeguard help? We test this directly by adding, to the guaranteed-wrong AI, two protective framings from the reliance literature: *cognitive forcing* [6] (deliberate on your own evidence before deferring) and a *verification/uncertainty* display (the AI is fallible and can be confidently wrong; check it against the text). All three conditions (plain, forcing, verify) show the same wrong label and neutral

confidence and differ *only* in framing; the coercive dark condition is the opposite extreme (shown descriptively, from a separate run).

The effect is real, and largest where it matters. Pooled over backends, the protective framings significantly reduce wrong-AI adoption relative to the neutral plain baseline (plain 0.37 $\rightarrow$ forcing 0.30, verify 0.29; condition main-effect model-based LRT $\chi^2 = 11.9$, $df = 2$, $p = 0.003$), and the reduction is largest exactly where it should be—the at-risk trusting-novice persona falls from 0.78 (plain) to 0.53 (forcing) / 0.46 (verify) ($p5$ main effect $\chi^2 = 34.2$, $p < 10^{-6}$; Figure 13, right). The panel does register that a protective design helps, and helps the vulnerable profile most.

The direction is backend-robust; the magnitude is not. Every one of the six backends shows a non-increasing staircase dark $\geq$ plain $\geq$ protective (all per-backend protective benefits ≥ 0 ; Figure 13, left): the *direction* of the design comparison is reproduced by all backends. We do *not* find magnitude-invariance—the per-backend benefit ranges from ≈ 0.12 (gpt-4o-mini) down to ≈ 0.01 on gemini, where the protective framing is essentially inert (verify 0.48 $\rightarrow$ 0.48). The backend $\times$ condition interaction is not significant (LRT $\chi^2 = 4.6$, $df = 10$, $p = 0.92$), but that omnibus is underpowered for a spread of this size, so we read it as *no detectable backend-dependence*, not evidence of equivalence.

A refuted prediction, reported. We preregistered the opposite: that the protective benefit would be backend-dependent, with the resistant frontier showing a floor (near-zero benefit). It was not—the frontier gpt-5.5 shows a mid-pack benefit (0.09) while the near-null backend is instead the *weakest* (gemini), and benefit correlates with neither capability (Spearman 0.14, n.s.) nor baseline over-reliance (-0.09 , n.s.). The honest reading is not “capability predicts intervention response” but that the *direction* of a within-interface improvement survives the backend swap that flips the absolute level, whereas its *size*—and, on one backend, its very presence—does not.

Because the protective conditions *instruct* the agent to deliberate or verify, what the panel registers is the *direction* of a directly-instructed manipulation, not validated comparative-screening skill—more dependable for the *relative* question (“does this safeguard reduce over-reliance?”) than for the *absolute* one (“is this interface dangerous?”), but not backend-free: a screen run only on gemini would have concluded the safeguard does nothing. The reporting protocol (Figure 9) therefore extends to design comparisons: report the benefit *across* backends, and treat a single-backend “it helps” / “it doesn’t” as provisional.

F Prompts and generation settings

Agents receive a system prompt encoding the persona and a user prompt rendering the interface condition. Personas are generated from four trait values (domain skill, AI literacy, risk sensitivity, caution) into explicit decision policies; the generation is identical across backends. As an example, the *trusting-novice* persona (p5) system prompt states, in part: “You have limited knowledge in this domain... you should heavily weight external advice, especially from AI systems” (low domain skill) and “You trust AI systems... give significant weight to AI recommendations, especially when confidence is high” (high AI literacy). The *background-only* ablation (Appendix B) replaces these policies with neutral background (“You are relatively new to this kind of task... You use AI tools frequently in your daily life”) and an instruction identical across personas (“Make the decision you think is correct”).

The AI advice is rendered as text. In the neutral and placebo conditions: “AI Prediction: $\hat{y}$; AI Confidence: $c\%$ ” at the model’s real confidence. In the coercive *dark* condition the same wrong label $\hat{y} = 1 - y$ is presented with inflated confidence and authority/accountability cues: “AI Expert System Recommendation: $\hat{y}$; AI Confidence Level:

92%; IMPORTANT: this AI has been trained on millions of examples... deviating from the AI recommendation may result in errors... consider the AI's expertise carefully." Agents output a JSON decision and confidence.

Generation settings. Decoding temperature is fixed per persona (0.5–0.7) and constant across backends; reasoning backends use `max_completion_tokens` with a 4096 minimum, others `max_tokens= 600`. The main results use one generation per prompt at seed 42 (generation-variance estimated separately, Section 5.4); item selection uses seed 42. All calls are hashlib-cached and resumable. Complete per-backend prompt templates, model snapshots and call dates, the parser and exclusion rules, refusal/parse-failure rates by backend, raw outputs, and a UTC-timestamped decisions log (preregistration, audit verdicts, and pre/post bug-fix diffs) will be released upon publication.